

CUDA by examples: delay and sum

Piotr Jarosik, Marcin Lewandowski, Billy Yiu

IUS 2023 GPU short course, Canada, Montreal

Agenda

Introduction

Delay

Sum

Linear array

Matrix array

Row Column Array

Motivation

Why delay and sum?

- ▶ GPU-based ***ultrasound*** signal processing, IUS 2023
- ▶ fundamental part of ultrasound signal processing
- ▶ combination of two more general problems: interpolation and reduction
- ▶ good practical CUDA example

Objectives


Practical introduction to:

- ▶ CUDA kernel syntax (grids, blocks, global, device kernels...)
- ▶ CUDA memory model (global, shared, texture)
- ▶ CUDA core performance guidelines

For more examples, see also:

<https://github.com/lab4us/gpu-short-course>

Delay and sum



Delay and sum


$$g(t, x, z) = \sum_{r=1}^{N_r} f(t, r, \delta(t, r, x, z)) \quad (1)$$

where:

- ▶ g : low resolution images (*LRI*)
- ▶ f : input signal (e.g. raw channel data)
- ▶ t : transmit index (e.g. plane wave transmit number)
- ▶ r : receiver index, N_r : total number of receivers
- ▶ x, z : output image x, z coordinates
- ▶ $\delta(t, r, x, z)$: beamforming delays

To get the final B-mode image: compound, envelope, logarithmic compression.

Plane wave imaging



$$\delta(t, r, x, z) = (z \cos(\alpha(t)) + x \sin(\alpha(t)) + \sqrt{(x - l(r))^2 + z^2}) \frac{f_s}{c}$$

where:

- ▶ $l(r)$ is the position of probe element r (*azimuth*),
- ▶ $\alpha(t)$ is an angle in transmit t

Delay and sum decomposition

How we interpret delay and sum?


- ▶ delay \equiv interpolation (signal *sampling* with delays)

$$g(t, r, x, z) \leftarrow f(t, r, \delta(t, r, x, z))$$

- ▶ sum \equiv just sum :) (an example of data *reduction*)

$$h(t, x, z) \leftarrow \sum_{r=1}^{N_r} g(t, r, x, z)$$

Delay




$$g(t, r, p) \leftarrow f(t, r, \delta(t, r, p))$$

where $p = (x, z)$.

Delay \equiv interpolation


- ▶ "Delay" is the data sampling part of processing.
- ▶ The signal is sampled according to the beamforming delays δ .



source: So you think you can DAS? Vincent Perrot et al., 2021


Delay \equiv interpolation

- ▶ $\delta : \rightarrow \mathbb{R}$, δ can be a floating-point value
- ▶ our raw channel data is a discretized signal (sampled with finite sampling frequency)
- ▶ we have to **interpolate** in case δ points somewhere between samples
 - ▶ it would be also good to **extrapolate** to reasonable value, e.g. 0



source: <https://commons.wikimedia.org/>

CUDA: interpolation kernel – block of threads



CUDA: 1D nearest-neighbor interpolation

```
__global__  
void interpolate1d(  
    float *output, float *input, float *points,  
    size_t np, size_t ns  
) {  
    int x = blockIdx.x*blockDim.x + threadIdx.x;  
  
    if(x >= np) {  
        return;  
    }  
  
    int point = int(roundf(points[x]));  
    if(point < ns) {  
        output[x] = input[point];  
    }  
    else {  
        return 0.0f;  
    }  
}
```

CUDA: 1D nearest-neighbor interpolation

```
def interpolate1d(input_array, points):
    interpolate1d_kernel = load_cuda_kernel("1_1_interpolate1d.cc")
    n_samples, n_points = len(input_array), len(points)

    block = (256, )
    grid = (math.ceil(n_points/256), )

    output_array = cp.zeros((n_points, ), dtype=np.float32)
    interpolate1d_kernel(
        grid, block,
        args=(
            output_array,
            cp.asarray(input_array).astype(cp.float32),
            cp.asarray(points).astype(cp.float32),
            n_points, n_samples
        )
    )
    return output_array
```

CUDA: 1D linear interpolation

```
__global__
void interpolate1d(
    float *out, float *in, float *points,
    size_t np, size_t ns
) {
    int x = blockIdx.x*blockDim.x + threadIdx.x;
    if(x >= np) {
        return;
    }

    float sample = points[x];
    float weight = modff(sample, &sample);

    int sample_nr = int(sample);
    if(sample_nr < ns-1) {
        return in[sample_nr]*(1-weight) + in[sample_nr+1]*weight;
    }
    else {
        return 0.0f;
    }
}
```

CUDA: 1D linear interpolation

```
__forceinline__
__device__
float interp_linear(float *in, float sample, int ns) {
    float weight = modff(sample, &sample);
    int sample_nr = int(sample);
    if(sample_nr < ns-1) {
        return in[sample_nr]*(1-weight) + in[sample_nr+1]*weight;
    }
    else {
        return 0.0f;
    }
}

__global__
void interpolate1d(float *out, float *in, float *points, size_t np, size_t ns) {
    int x = blockIdx.x*blockDim.x + threadIdx.x;
    if(x >= np) {
        return;
    }
    out[x] = interp_linear(in, points[x], ns);
}
```

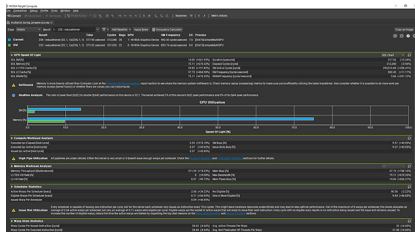


Profiler

How can we measure kernel's performance?

We can analyse time and memory requirements using **code profiler**.

NVIDIA provides two profiling tools:

- ▶ Nsight Systems: overall performance of the CUDA streams
- ▶ Nsight Compute: detailed analysis of a CUDA kernel



CUDA: 1D linear interpolation

```
nsys profile --stats=true --trace cuda python 1_1_interpolate1d.py
```

NVIDIA Geforce 3700 Mobile NVIDIA GPU 3700 RTX
Results for 2^{15} samples:

```
...
Generating '/tmp/nsys-report-425c.qdstrm'
[1/6] [=====100%] report1.nsys-rep
[2/6] [=====100%] report1.sqlite
...
[4/6] Executing 'cuda_gpu_kern_sum' stats report
```

Time (%)	Total Time (ns)	Instances	Avg (ns)	Med (ns)	Min (ns)	Max (ns)	StdDev (ns)	Name
49,0	184 673	10	18 467,0	18 464,0	18 336	18 624	109,0	interpolate1d

CUDA: 1D linear interpolation for DAS

Now let's move to the *layered* version of our kernel.

We will use it in DAS implementation.

The *layered* kernel should perform 1D interpolation for each t , r and p in parallel.

$$g(t, r, p) \leftarrow f(t, r, \delta(t, r, p))$$

CUDA: 1D linear interpolation – DAS

Input array: $(N_t, N_r, N_s) = (3, 128, 8192)$

Output array: $(N_t, N_r, N_p) = (3, 128, 256 * 1024)$

Thread block: $(R_x, T_x, Points)$

CUDA: 2D linear interpolation – first attempt

```
__global__  
void interpolate1d_layered(  
    float *out, float *in, float *points,  
    size_t nt, size_t ne, size_t ns, size_t np  
) {  
    int e = blockIdx.x*blockDim.x + threadIdx.x;  
    int t = blockIdx.y*blockDim.y + threadIdx.y;  
    int p = blockIdx.z*blockDim.z + threadIdx.z;  
  
    if(p >= np || e >= ne || t >= nt) {  
        return;  
    }  
  
    // input [t, e, 0]  
    size_t in_i = (ne*ns)*t + (ns)*e;  
    // output [t, e, p]  
    size_t out_i = (ne*np)*t + (np)*e + p;  
  
    out[out_i] = interp_linear(&in[in_i], points[out_i], ns);  
}
```

CUDA: 2D linear interpolation – first-attempt

Step	Time (median, ms)
1D interpolation - first attempt	147

CUDA: 2D linear interpolation – first-attempt


Let's check global memory store efficiency:

```
ncu
--target-processes all
--metrics
    smsp__sass_average_data_bytes_per_sector_mem_global_op_st
-f python 1_2_interpolate1d_layered_1st_attempt.py
```


```
interpolate1d_layered (16, 1, 32768)x(8, 8, 8), Context 1, Stream 7, Device 0, CC 8.6
Section: Command line profiler metrics
```

Metric Name	Metric Unit	Metric Value
smsp__sass_average_data_bytes_per_sector_mem_global_op_st.max_rate	byte/sector	32
smsp__sass_average_data_bytes_per_sector_mem_global_op_st.pct	%	12,50
smsp__sass_average_data_bytes_per_sector_mem_global_op_st.ratio	byte/sector	4

Memory coalescing



Memory coalescing



CUDA: 1D linear interpolation – write coalescing

From:

```
int e = blockIdx.x*blockDim.x + threadIdx.x;  
int t = blockIdx.y*blockDim.y + threadIdx.y;  
int p = blockIdx.z*blockDim.z + threadIdx.z;
```

To:

```
int p = blockIdx.x*blockDim.x + threadIdx.x;  
int e = blockIdx.y*blockDim.y + threadIdx.y;  
int t = blockIdx.z*blockDim.z + threadIdx.z;
```

And Thread block: ($Points$, R_x , T_x)

CUDA: 2D linear interpolation – write-coalescing

Step	Time (median, ms)
1D interpolation - first attempt	147
1D interpolation - write coalescing	51

CUDA: 2D linear interpolation – write-coalescing

```
ncu
--target-processes all
--metrics
    smsp__sass_average_data_bytes_per_sector_mem_global_op_st
-f python 1_3_interpolate1d_layered_improved.py
```

```
interpolate1d_layered (32768, 16, 1)x(8, 8, 8), Context 1, Stream 7, Device 0, CC 8.6
Section: Command line profiler metrics
```

Metric Name	Metric Unit	Metric Value
smsp__sass_average_data_bytes_per_sector_mem_global_op_st.max_rate	byte/sector	32
smsp__sass_average_data_bytes_per_sector_mem_global_op_st.pct	%	100
smsp__sass_average_data_bytes_per_sector_mem_global_op_st.ratio	byte/sector	32

CUDA: 2D linear interpolation – compute-bound

For each output point interpolation we need to read one value from the `*delays` array.

Can we calculate delays on-the fly instead? Will the processing time be improved?

- ▶ the kernel becomes specific for some imaging approach
- ▶ however the kernel becomes *less memory-bound* and *more compute-bound*

CUDA: 1D linear interpolation – compute-bound

From:

```
float sample = points[out_i]
```

To:

```
int point_x = point/nz;
```

```
int point_z = point-point_x*nz;
```

```
float angle = alpha[tx];
```

```
float z = z_origin + point_z*dz;
```

```
float x = x_origin + point_x*dx;
```

```
float element_pos = probe_origin + element*pitch;
```

```
float tx_delay = (z*cos(angle) + x*sin(angle))/c;
```

```
float rx_delay = hypotf(z, (x-element_pos))/c;
```

```
float sample = (tx_delay+rx_delay)*fs;
```

CUDA: 2D linear interpolation – compute-bound


Step	Time (median, ms)
1D interpolation - first attempt	147
1D interpolation - write coalescing	51
1D interpolation - compute-bound	28

Texture memory

Texture memory is a GPU-specific global memory cache.

- ▶ CUDA-available cache.
- ▶ Dedicated for global memory access patterns with *spatial proximity*.
- ▶ Provides some basic interpolation and extrapolation functionality.

Texture memory



Texture memory

CUDA texture memory functionality:

- ▶ interpolation:
 - ▶ nearest-neighbor
 - ▶ linear, bilinear and trilinear
- ▶ extrapolation: clamp (extrapolate with the values on the border), constant value, ...
- ▶ address normalization (to $[0, 1]$)
- ▶ ...

CUDA: linear interpolation – Texture memory

```
from cp.cuda.texture import *

ch = ChannelFormatDescriptor(32, 0, 0, 0, cudaChannelFormatKindFloat)
array = CUDAarray(
    ch,
    width=n_samples,
    height=n_elements,
    depth=n_tx,
    flags=1, # CudaArrayLayered
)
res = ResourceDescriptor(cudaResourceTypeArray, cuArr=array)
tex = TextureDescriptor(
    (cudaAddressModeBorder, cudaAddressModeBorder),
    cudaFilterModeLinear,
    cudaReadModeElementType,
    borderColors=(0.0, 0.0, 0.0, 0.0)
)
texobj = TextureObject(res, tex)
# ... and instead of input array:
kernel(grid, block, args=(out, texobj, ...))
```

CUDA: linear interpolation – Texture memory

From:

```
out[out_i] = interp_linear(&in[in_i], points[out_i], ns);
```

To:

```
out[out_i] = tex2DLayered<float>(input, sample+0.5, element+0.5, t);
```

CUDA: 2D linear interpolation with Texture Memory

Step	Time (median, ms)
1D interpolation - first attempt	147
1D interpolation - write coalescing	51
1D interpolation - compute-bound	28
1D interpolation - texture memory	23

CUDA: 2D linear interpolation with Texture Memory

```
__global__
void interpolate1d_layered(
    float *output, cudaTextureObject_t input, const float *alpha,
    const size_t n_tx, const size_t n_elements, const size_t n_samples,
    const size_t n_points,
    const float c, const float fs,
    const float pitch, const float probe_origin,
    const float dx,    const float nx, const float x_origin,
    const float dz,    const float nz, const float z_origin) {

    int point = blockIdx.x*blockDim.x + threadIdx.x;
    int element = blockIdx.y*blockDim.y + threadIdx.y;
    int tx = blockIdx.z*blockDim.z + threadIdx.z;


    if(point >= n_points || element >= n_elements || tx >= n_tx) {
        return;
    }
    int point_x = point/nz;
    int point_z = point-point_x*nz;

    float angle = alpha[tx];
    float z = z_origin + point_z*dz;
    float x = x_origin + point_x*dx;
    float element_pos = probe_origin + element*pitch;

    float tx_delay = (z*cos(angle) + x*sin(angle))/c;
    float rx_delay = hypotf(z, (x-element_pos))/c;
    float sample = (tx_delay+rx_delay)*fs;


    size_t output_idx = (n_elements*n_points)*tx + (n_points)*element + point;
    output[output_idx] = tex2DLayered<float>(input, sample+0.5, element+0.5, tx);
}
```

Sum



$$h(t, p) \leftarrow \sum_{r=1}^R g(t, r, p)$$

Sequential sum – 1D



Sequential sum - 1D

```
__global__  
void sum_sequential(float* output, float *input, size_t ne) {  
    float result = 0.0f;  
    for (int e = 0; e < ne; ++e) {  
        result += input[e];  
    }  
    output[0] = result;  
}
```

Sequential sum - 1D

Block and grid: (1,)

Vector of 1024 elements.

Step	Time (median, μs)
Sequential sum	340

Sequential sum - 1D

- ▶ $O(n)$
- ▶ optimal for a single-threaded implementation
- ▶ but only a single thread is used, and GPU have them a lot more

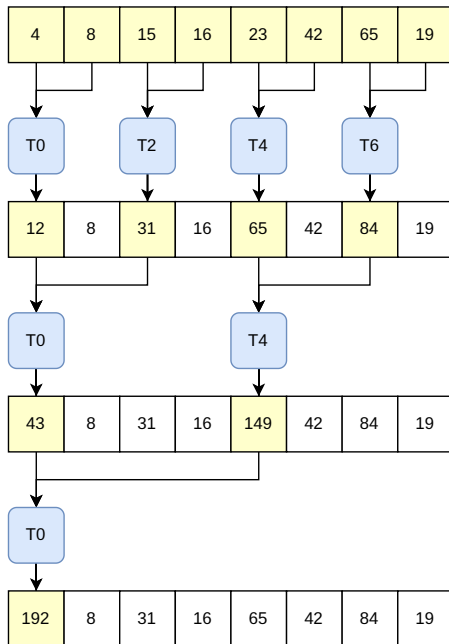
Parallel sum - 1D

We can further reduce the computational complexity to $O(\lg n)$ for the multi-core CPU/GPUs:

- ▶ threads with id divisible by 2: $h[id]_+ = h[id + 1]$ then,
- ▶ threads with id divisible by 4: $h[id]_+ = h[id + 2]$ then,
- ▶ ...
- ▶ threads with id divisible by $2^{\lg(n)}$: $h[id]_+ = h[id + 2^{n-1}]$.


This is Brent's algorithm, usually visualized as a reduction tree.

Parallel sum - 1D



Parallel sum - 1D

- ▶ The reduction can be performed by a single block of threads, using shared memory.
- ▶ However the max number of threads ≤ 1024 .
- ▶ How can we reduce larger arrays? Apply the reduction tree recursively.



source: Optimizing Parallel Reduction in CUDA, Mark Harris

Parallel sum - 1D

```
__global__
void sum_parallel(float *out, float *in) {
    extern __shared__ float shared_memory[];
    int thread = threadIdx.x;
    int element = thread;

    shared_memory[thread] = in[element];
    __syncthreads();
    for(int s = 1; s < blockDim.x; s *= 2) {
        if (thread % (2*s) == 0) {
            shared_memory[thread] += shared_memory[thread + s];
        }
        __syncthreads();
    }
    if (thread == 0) {
        out[nz*x+z] = shared_memory[0];
    }
}
```

Parallel sum - 1D

Block and grid: (1024,)

Vector of 1024 elements.

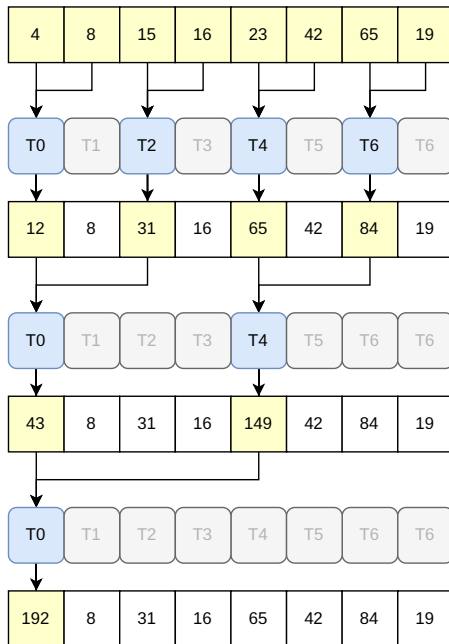
Step	Time (median, μs)
Sequential sum	340
Parallel sum	37

Parallel sum - 1D – branch divergence

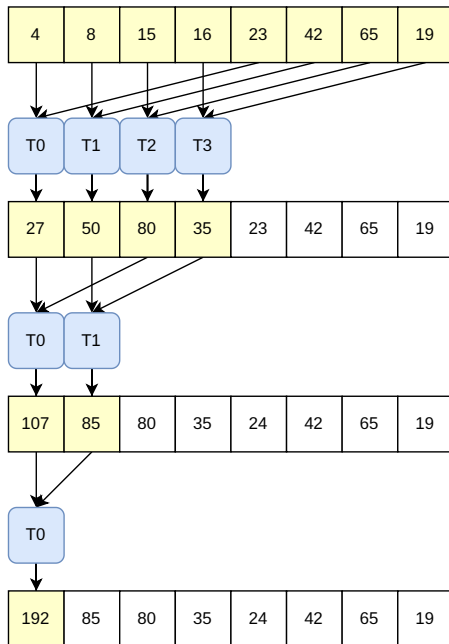
Every second thread in the warp (32 threads) is inactive.

```
for(int s = 1; s < blockDim.x; s *= 2) {  
    if (thread % (2*s) == 0) {  
        shared_memory[thread] += shared_memory[thread+s];  
    }  
    __syncthreads();  
}
```

Parallel sum - 1D – branch divergence



Parallel sum - 1D – non-divergent



Parallel sum - 1D – non-divergent

From:

```
for(int s = 1; s < blockDim.x; s *= 2) {  
    if (thread % (2*s) == 0) {  
        shared_memory[thread] += shared_memory[thread+s];  
    }  
    __syncthreads();  
}
```

To:

```
for (int s = blockDim.x/2; s > 0; s >>= 1) {  
    if (thread < s) {  
        shared_memory[thread] += shared_memory[thread+s];  
    }  
    __syncthreads();  
}
```

Parallel sum – 1D – non-divergent

Step	Time (median, μs)
Sequential sum	340
Parallel sum	37
Parallel sum – non-divergent	23

Parallel sum – 1D – non-divergent

Only half of the threads is active on the first loop (and the other half is just inactive).

- ▶ Halve the number of threads/blocks.
- ▶ Replace single element load with two loads and a single addition.

Parallel sum – 1D – non-divergent

From:

```
shared_memory[thread] = in[element];
```

To:

```
shared_memory[thread] = in[element]+in[element+blockDim.x];
```

Parallel sum – 1D – first add

Block and grid: (1024,)

Vector of 1024 elements.

Step	Time (median, μs)
Sequential sum	340
Parallel sum	37
Parallel sum – non-divergent	22
Parallel sum – first add	16

Parallel sum – 1D – unroll last warp

- ▶ when $s \leq 32$ only one warp in the block is active
- ▶ SIMD
 - ▶ we don't need `syncthreads`
 - ▶ we can skip check if thread id is $< s$ (all threads within warp are used anyway)

Parallel sum – 1D – unroll last warp

From:

```
for (int s = blockDim.x/2; s > 0; s >>= 1) {  
    if(thread < s) {  
        shared_memory[thread] += shared_memory[thread+s];  
    }  
    __syncthreads();  
}
```

To:

```
for (int s = blockDim.x/2; s > 32; s >>= 1) {  
    if(thread < s) {  
        shared_memory[thread] += shared_memory[thread+s];  
    }  
    __syncthreads();  
}  
if(thread < 32) {  
    volatile float* sm = shared_memory;  
    sm[thread] += sm[thread + 32];  
    sm[thread] += sm[thread + 16];  
    sm[thread] += sm[thread + 8];  
    sm[thread] += sm[thread + 4];  
    sm[thread] += sm[thread + 2];  
    sm[thread] += sm[thread + 1];  
}
```

Parallel sum – 1D – unroll last warp


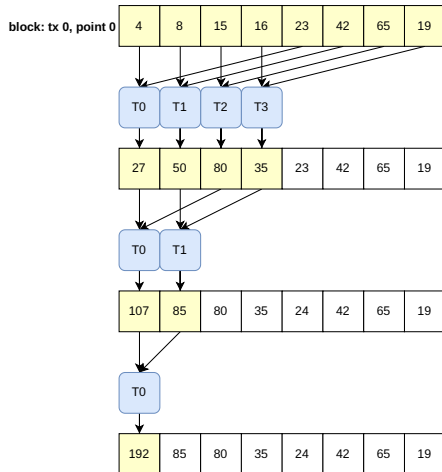
Block and grid: (1024,)

Vector of 1024 elements.

Step	Time (median, μs)
Sequential sum	340
Parallel sum	37
Parallel sum – non-divergent	22
Parallel sum – first add	16
Parallel sum – unroll last warp	12

Parallel sum – DAS

How can we assign reduction for each output image point?



Parallel sum – DAS

```
__global__
void sum_parallel4(float *out, float *in,
                  size_t nx, size_t nz, size_t n_elements)
{
    extern __shared__ float shared_memory[];
    int thread = threadIdx.x;
    int z = blockIdx.y;
    int x = blockIdx.z;
    int el = thread;
    if (z >= nz || x >= nx) {return;}

    int off = nz*n_elements*x + n_elements*z;
    shared_memory[thread] = in[off+el] + in[off+el+blockDim.x];
    __syncthreads();

    for (int s = blockDim.x/2; s > 32; s >>= 1) {
        if (thread < s) {
            shared_memory[thread] += shared_memory[thread+s];
        }
        __syncthreads();
    }
    if(thread < 32) {
        volatile float* warp_shm = shared_memory;
        warp_shm[thread] += warp_shm[thread + 32];
        warp_shm[thread] += warp_shm[thread + 16];
        warp_shm[thread] += warp_shm[thread + 8];
        warp_shm[thread] += warp_shm[thread + 4];
        warp_shm[thread] += warp_shm[thread + 2];
        warp_shm[thread] += warp_shm[thread + 1];
    }

    if (thread == 0) {
        out[nz*x+z] = shared_memory[0];
    }
}
```

Parallel sum – DAS

Array (256, 1024, 128).

Step	Time (median, ms)
Sequential sum	7.87
Parallel sum	11.38
Parallel sum – non-divergent	7.82
Parallel sum – conflict free	7.71
Parallel sum – first add	7.63
Parallel sum – unroll last warp	7.62

Sum – Conclusions

Where is the bottleneck?

We have $O(\lg n)$, however the constant factor may be quite high:

- ▶ synchronization barrier,
- ▶ occupancy (with every loop iteration less and less threads are active),
- ▶ instruction overhead (see Mark Harris presentation on it).

And the axis we are reducing along has only 128 elements.

What options do we have?

- ▶ Use sequential sum: it's probably good enough for < 1024 elements.
- ▶ Combine sequential with parallel sum

In the further examples we will use sequential sum, for the sake of simplicity.

Combining delay and sum: stream of 2 kernels

Let's combine the kernels to get the reconstructed ultrasound images.

```
with stream:  
    h = interpolate1d_kernel(block1, grid1, params=(f, # ...  
    g = sequential_sum(block2, grid2, params=(h, #...
```


Combining delay and sum: stream of 2 kernels

Step	Time (median, ms)
Delay and sum: 2 kernels	$27+25 = 52$

Frame rate < 20 FPS.

Not impressive...

Combining delay and sum: stream of 2 kernels

Where is the bottleneck?

$$g(t, r, x, z) \leftarrow f(t, r, \delta(t, r, x, z))$$

$$h(t, x, z) \leftarrow \sum_{r=1}^R g(t, r, x, z)$$

- ▶ interp1d writes h to global memory,
- ▶ sum reads h from global memory,
- ▶ h has $N_t N_p N_r$ values,
- ▶ reading/writing from/to global memory is expensive.

Do we really need g ?

Answer: usually no, and we can combine delay and sum, e.g. implement kernel $h(f)$.

Combining delay and sum: single kernel

```
__global__
void beamform(float *out, const float *in, const float *angles,
              const float pitch, const float probe_origin, /*...*/) {
    int z = blockIdx.x*blockDim.x + threadIdx.x;
    int x = blockIdx.y*blockDim.y + threadIdx.y;
    int tx = blockIdx.z*blockDim.z + threadIdx.z;

    float rx_distance, sample_number, element_pos;
    int offset;

    if (z >= nz || x >= nx || tx >= n_angles) {
        return;
    }

    float pixel_x = x_origin + dx*x; // [m]
    float pixel_z = z_origin + dz*z; // [m]
    float angle = angles[tx];
    int tx_offset = (n_elements*n_samples)*tx;

    float tx_distance = pixel_z*cosf(angle) + pixel_x*sinf(angle);

    out[(nz*nx)*tx + (nz)*x + z] = 0.0f;
    for (int element = 0; element < n_elements; ++element) {
        offset = tx_offset + (n_samples)*element;

        element_pos = probe_orig + element*pitch;
        rx_distance = hypotf(pixel_x - element_pos, pixel_z);
        delay = (tx_distance + rx_distance)/speed_of_sound*sampling_frequency;

        float value = interp_linear(&in[offset], delay, n_samples);
        out[(nz*nx)*tx + (nz)*x + z] += value;
    }
}
```

Combining delay and sum: single kernel

Step	Time (median, ms)
Delay and sum: 2 kernels	$27+25 = 52$
Delay and sum: 1 kernel	8

Conclusions and further improvements

Conclusion: optimizing the memory access pattern is usually critical for performance.

Further improvements:

- ▶ use IQ data (and increase the decimation factor)
- ▶ combine sequential with parallel reduction
- ▶ use intrinsic functions (e.g. `sincos` to get `sin` and `cos`)
- ▶ move transmit parameters (e.g. transmit angle) to constant memory (in case $\gg 3$)

Delay and sum for matrix array

$$g(t, x, y, z) = \sum_{r_y=1}^{N_{r_y}} \sum_{r_x=1}^{N_{r_x}} f(t, r_x, r_y, \delta(t, r_x, r_y, x, y, z)) \quad (2)$$


Where *delta* depends on a pair of angles α, γ or OXZ, OYZ in case of PW imaging.

Some practical notes:



- ▶ Delays becomes $\delta(t, r_x, r_y, x, y, z)$ and can be large for extensive volumes – highly recommended to **compute on the fly** to save GPU RAM and reduce memory W/R.
- ▶ The LRIs become $g(t, x, y, z)$ and can be relatively large for extensive volumes – consider summing along TX and RX in a single kernel.
- ▶ δ becomes more complex due to 2D angle space – highly recommended to optimize with CUDA intrinsics.

Delay and sum for RCA

RCA – Row-column Arrays




Delay and sum for RCA



$$\|\vec{t}\|_2 = z \cos(\varphi) + x \sin(\varphi), \|\vec{r}_i\|_2 = \sqrt{z^2 + (y - r_i)^2}$$

$$\delta_t(t, x, z) = \|\vec{t}\|_2 \frac{fs}{c}, \delta_r(r, y, z) = \|\vec{r}_i\|_2 \frac{fs}{c}$$

Delay and sum for RCA



- ▶ D4L-05, Thursday, September 7, 2:45, *Real-Time 3D Imaging Pipeline for Row-Column Array Probe and Holographic Display*, Marcin Lewandowski et al.
- ▶ exhibition boot

References

- ▶ *Programming Massively Parallel Processors*, David Kirk, Wen-Mei W. Hwu
- ▶ *Optimizing Parallel Reduction in CUDA*, Mark Harris
- ▶ *CUDA Toolkit Documentation 12.2*, NVIDIA

CUDA by example: delay and sum

Thank you!

→ Billy, Real-time Speed of Sound Estimation using CNN
Inferencing Using GPU