

GPU-Based Ultrasound Signal Processing

Real-Time Speed of Sound Estimation
and CNN Inferencing Using GPU

Presenter: Dr Billy Yiu



Outline of Discussion

Execute your algorithm on GPU using existing libraries

Part I Real-time speed of sound estimation using GPU

- How can we use CuPy to execute the algorithm in real-time/close to real-time

Part II C++ Deployment of Machine Learning Algorithms

- How can we use ONNX and TensorRT to deploy a machine learning algorithm in a standalone program?


Real-Time Speed of Sound Estimation Using GPU

Di Xiao, Hassan Nahas, Billy Yiu


Ultrasound Imaging with Assumed Speed-of-Sound (SoS) Value

Challenge: Loss of image quality when beamformed with wrong SoS



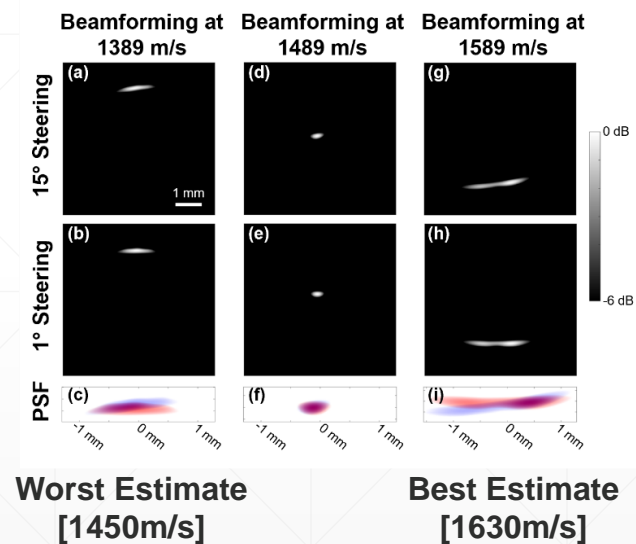
Plane Wave Imaging Forward Model



Effect of Different Beamforming SoS Values



Point Spread Function of Wire Target in Water



Global SoS Estimation: Modified Delay-and-Sum Beamforming


Goal: Provide a real-time estimate of the average (global) SoS of the imaged medium

1. Acquisition




Plane Wave RF Data
from N_A Angles


2. Speed-of-Sound Aware Beamforming



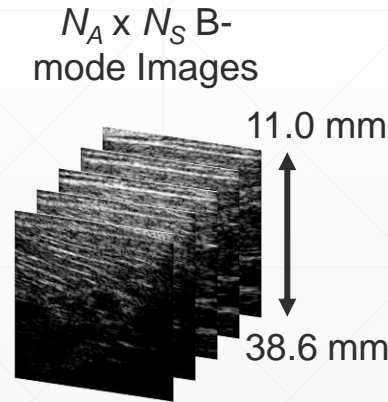
1400 m/s



⋮





1700 m/s



Beamform for a range of SoS values

Refraction Compensation




$$\theta_{\text{medium}} = \sin^{-1} \left(\sin \theta_{\text{probe}} \frac{v_{\text{medium}}}{v_{\text{assumed}}} \right)$$

Snell's Law


Global SoS Estimation: Identifying Best Beamforming SoS

Goal: Provide a real-time estimate of the average (global) SoS of the imaged medium

3. Inter-steering Loss




Inter-Steering Loss Calculation



Calculate a loss for each SoS:
differences between steering angles

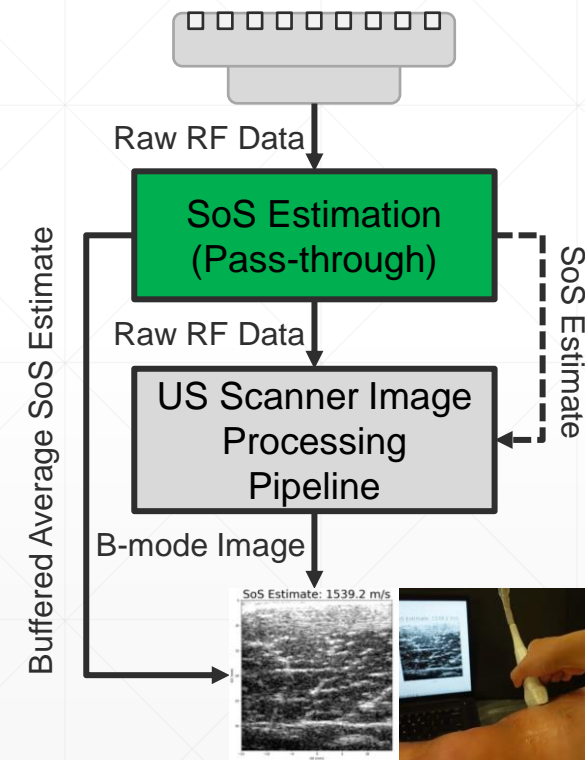
4. SoS Estimation

In vivo Example Loss





Estimate the SoS by finding
the smoothed loss minimum

Real-Time Global SoS Estimation



Display the SoS estimate in
real-time alongside B-mode

A High-Level Overview of the GPU Implementation of the SOS Estimation Algorithm




Python Code Structure for Real-Time SoS Estimation

Leverage the customizable receive processing pipeline

```
scheme = Scheme(  
    tx_rx_sequence=sequence,  
    processing=Pipeline(  
        steps=(  
            Lambda(lambda data: (data_timestamps.append(time_ns()), data)[1]),  
            RemapToLogicalOrder(),  
            Transpose(axes=(2, 3, 1, 0)),  
            Lambda(lambda data: (data_buffer_rf.append(data.get()), data[:, :, :Nang, :])[1]),  
            Lambda(pass_through_bmfrm, lambda metadata: metadata.copy(input_shape=(Nz, Nx), dtype="float32")),  
        ),  
        placement="/GPU:0"  
    ))
```

Pipeline: Automatically execute the specified functions sequentially

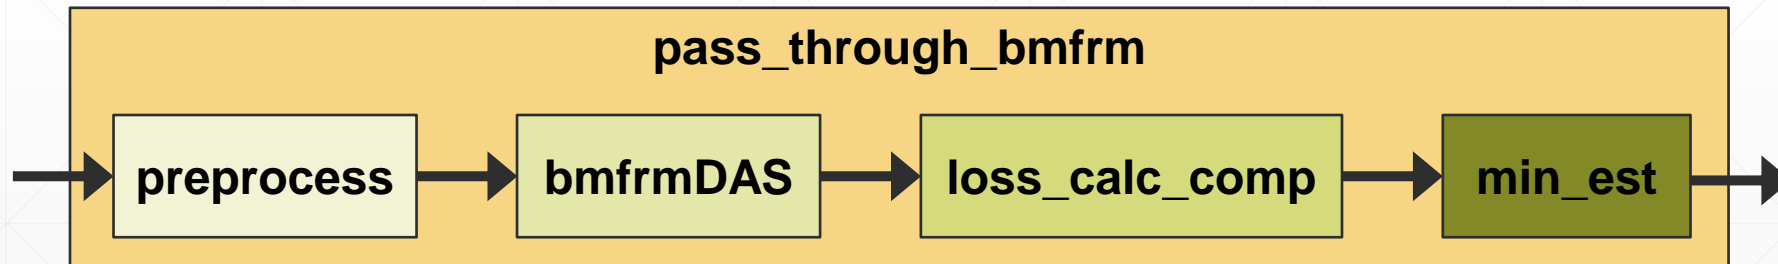


Lambda function: A function handler that the pipeline can call

Inside the Pass_Through_Bmfrm...

```
def pass_through_bmfrm(data):  
    data_preprocess = preprocess(data)  
    b_img = bmfrmDAS(data_preprocess) ← sos aware BF  
    sos_est = min_est(loss_calc_comp(data_preprocess).get(), sosBmfrm, sosTest)  
    data_buffer_sos.append(sos_est)  
    print("FPS: {:.31f} fps".format(fps_calc(data_timestamps)), end='\r')  
    return b_img
```

Modularize the processing pipeline



These modules can be CuPy built-in or custom functions

Inside the Pass_Through_Bmfrm...

```
def pass_through_bmfrm(data):  
    data_preprocess = preprocess(data)  
    b_img = bmfrmDAS(data_preprocess)  
    sos_est = min_est(loss_calc_comp(data_preprocess).get(), sosBmfrm, sosTest)  
    data_buffer_sos.append(sos_est)  
    print("FPS: {:.31f} fps".format(fps_calc(data_timestamps)), end='\r')  
    return b_img
```

These modules can be CuPy built-in or custom functions

- **Divide the overall process into multiple steps**
 - Allows optimized implementation for each step
 - All steps are implemented using CuPy functions
 - Some of the functions can be implemented using custom kernels
- **Same as C++ implementation**
 - Use pre-built libraries for convenience
 - Built custom implementation for mission-critical steps

Implementation of the PreProcess Module

```
def preprocess(rfTensor):  
    data_preprocess = hilbert_fft(bandpass_filter(rfTensor.astype('float32')))  
    data_preprocess = cp.reshape(data_preprocess, (Nsamp*Nchan*Nang,), order='F')  
    return data_preprocess
```

bandpass_filter: Suppress out-of-band noise

```
def bandpass_filter(rfTensor):  
    return cp.flip(convolve1d(cp.flip(convolve1d(rfTensor, cp.squeeze(filtcoeff_gpu), axis = 0), axis = 0), cp.squeeze(filtcoeff_gpu), axis = 0), axis = 0)
```

Perform forward-backward FIR filtering using convolve1d from cupyx.scipy

hilbert_fft: Convert the RF data into its analytic form

```
def hilbert_fft(rfTensor):  
    rf_fft = cufft.fft(rfTensor, axis = 0)  
    rf_fft[(rfTensor.shape[0]/2):, :, :] = 0  
    return cufft.ifft(rf_fft, axis = 0)
```

Perform analytic signal conversion using Hilbert transform (zero-out the -ve side of the frequency spectrum)

reshape: reshape the filtered data into a column vector format for the next step

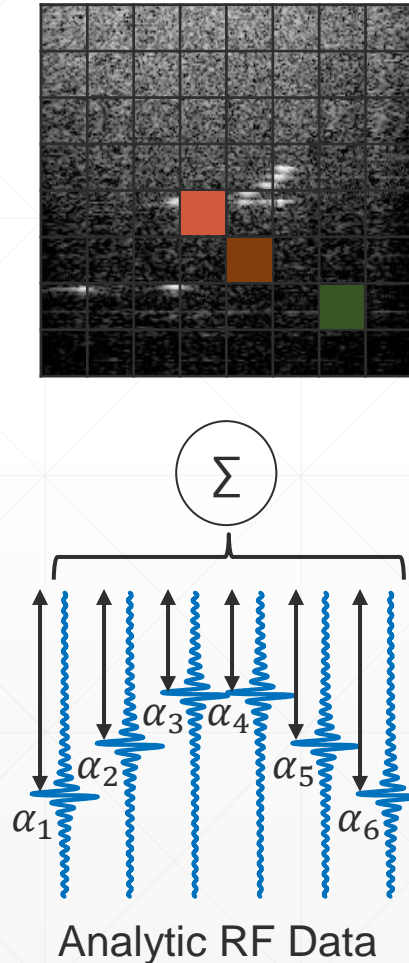
Implementation of the Beamforming Module

```
def bmfrmDAS(data_preprocess):  
    img=cp.reshape(sp_csr_bmfrm.dot(data_preprocess),(Nz*Nx,Nang),order='F')  
    img[img==0] = cp.asarray(np.nan*0j)  
    img = 20*cp.log10(cp.abs(cp.nanmean(img,axis=1)))  
    img = cp.nan_to_num(img, copy=False, nan=-100)  
    img = cp.reshape(img,(Nz,Nx),order='F')  
    return img
```

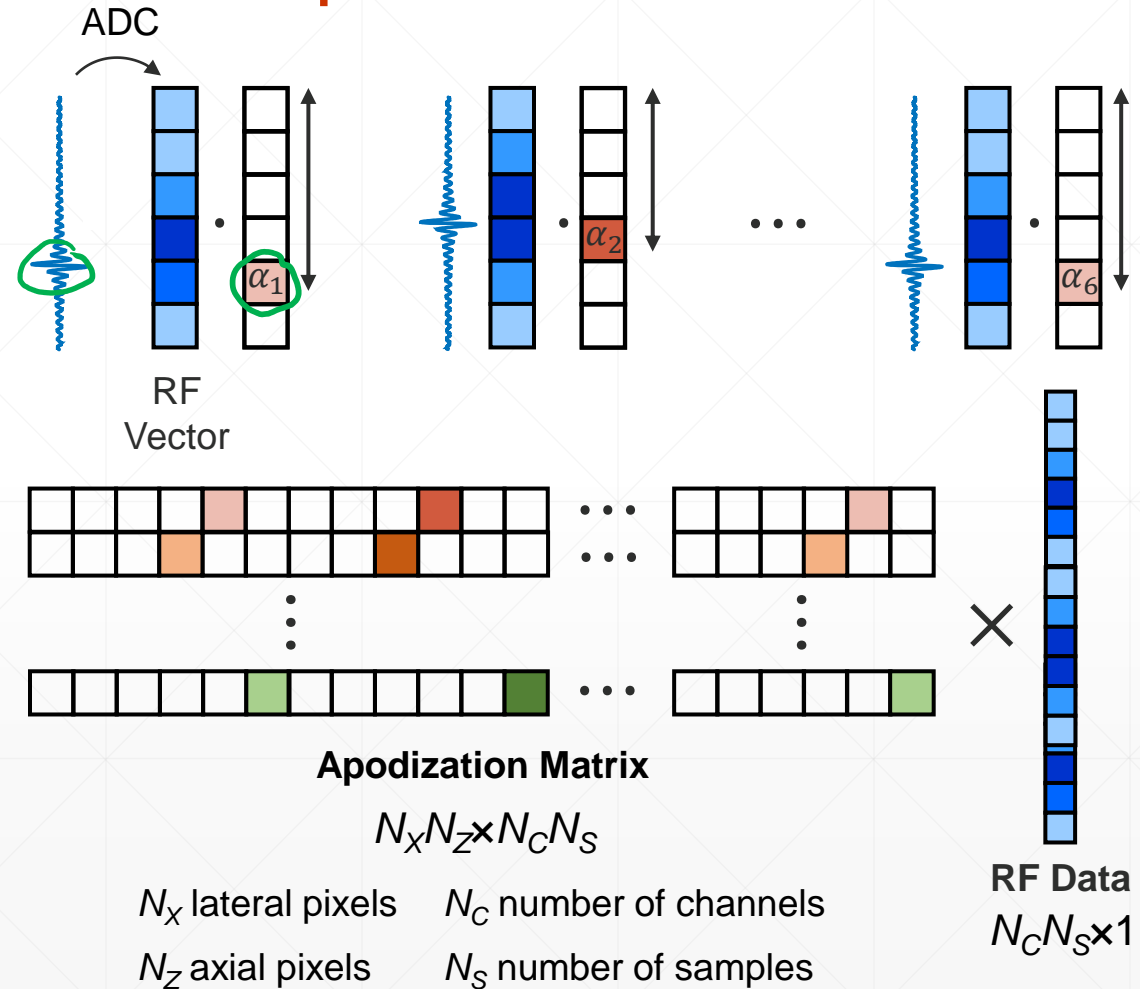
**Sparse matrix beamforming – Multiply the RF data matrix
with pre-computed delay matrices**

Brief Summary of Sparse Matrix Beamforming

Pixel by Pixel Approach



Expression in Matricial Form



Implementation of the Beamforming Module

```
def bmfrmDAS(data_preprocess):  
    img=cp.reshape(sp_csr_bmfrm.dot(data_preprocess),(Nz*Nx,Nang),order='F')  
    img[img==0] = cp.asarray(np.nan*0j)  
    img = 20*cp.log10(cp.abs(cp.nanmean(img,axis=1)))  
    img = cp.nan_to_num(img, copy=False, nan=-100)  
    img = cp.reshape(img,(Nz,Nx),order='F')  
    return img
```

Sparse matrix beamforming – Multiply the RF data matrix with pre-computed delay matrices

1. **Perform matrix multiplication (dot product) between the delay-encoded sparse matrix with the RF data column vectors**
 - `sp_csr_bmfrm` is an CuPy sparse matrix
 - Apodization is also encoded into the sparse matrix as well
2. **Log compress the output vector and reshape it into a 2D matrix (i.e., an image)**

Easy and straightforward to implement

Implementation of Loss Function Computation

```
def loss_calc_comp(data_preprocess):  
    bimgs = cp.reshape(sp_csr.dot(data_preprocess), (Npix*Npix*Nsos, Nang), order='F')  
    bimgs[bimgs==0] = cp.asarray(np.nan*0j)  
    std_comp = cp.sqrt(cp.square(cp.nanstd(cp.real(bimgs), axis=1, ddof=1))  
        + cp.square(cp.nanstd(cp.imag(bimgs), axis=1, ddof=1)))  
    coefvar = cp.divide(std_comp, cp.nanmean(cp.abs(bimgs), axis=1))  
    → loss = cp.nansum(cp.reshape(coefvar, (Npix*Npix, Nsos), order='F'), axis=0)  
    return cp.divide(cp.squeeze(loss), norm_factor)
```

1. **Use sparse matrix beamforming to generate images with different SoS**
 - `sp_csr` is an CuPy sparse matrix that accounts for SoS when encoding the delay
2. **Compute the coefficient of variation (CoV) across the angles**
3. **Cumulate the CoV across the image as the loss for different pre-defined SoS**

Estimating the SoS from the Loss Function

```
def min_est(loss_est, sosBmfrm, sosTest):  
    spl = UnivariateSpline(sosBmfrm, loss_est, s=0)  
    int_points = spl(sosTest)  
    min_sos = sosTest[np.argmin(int_points)]  
    return min_sos
```

1. Interpolate the loss function with univariate spline function
 - Done in the CPU
2. Find and return the SoS at minimum

These modules can be CuPy built-in functions or custom functions

Comparison Between Python and C++ Implementation

Python Implementation

CuPy makes GPU implementation of your algorithms much easier

- Similar to MatLab script
- No need to implement everything by yourselves
- Speed up your algorithms for a real-time /close-to-real-time performance
- Based on cuBLAS, cuRAND, cuSOLVER, cuSPARSE, cuFFT, cuDNN
- Easy display options

C++ Implementation

C++ provides a standalone program for the execution

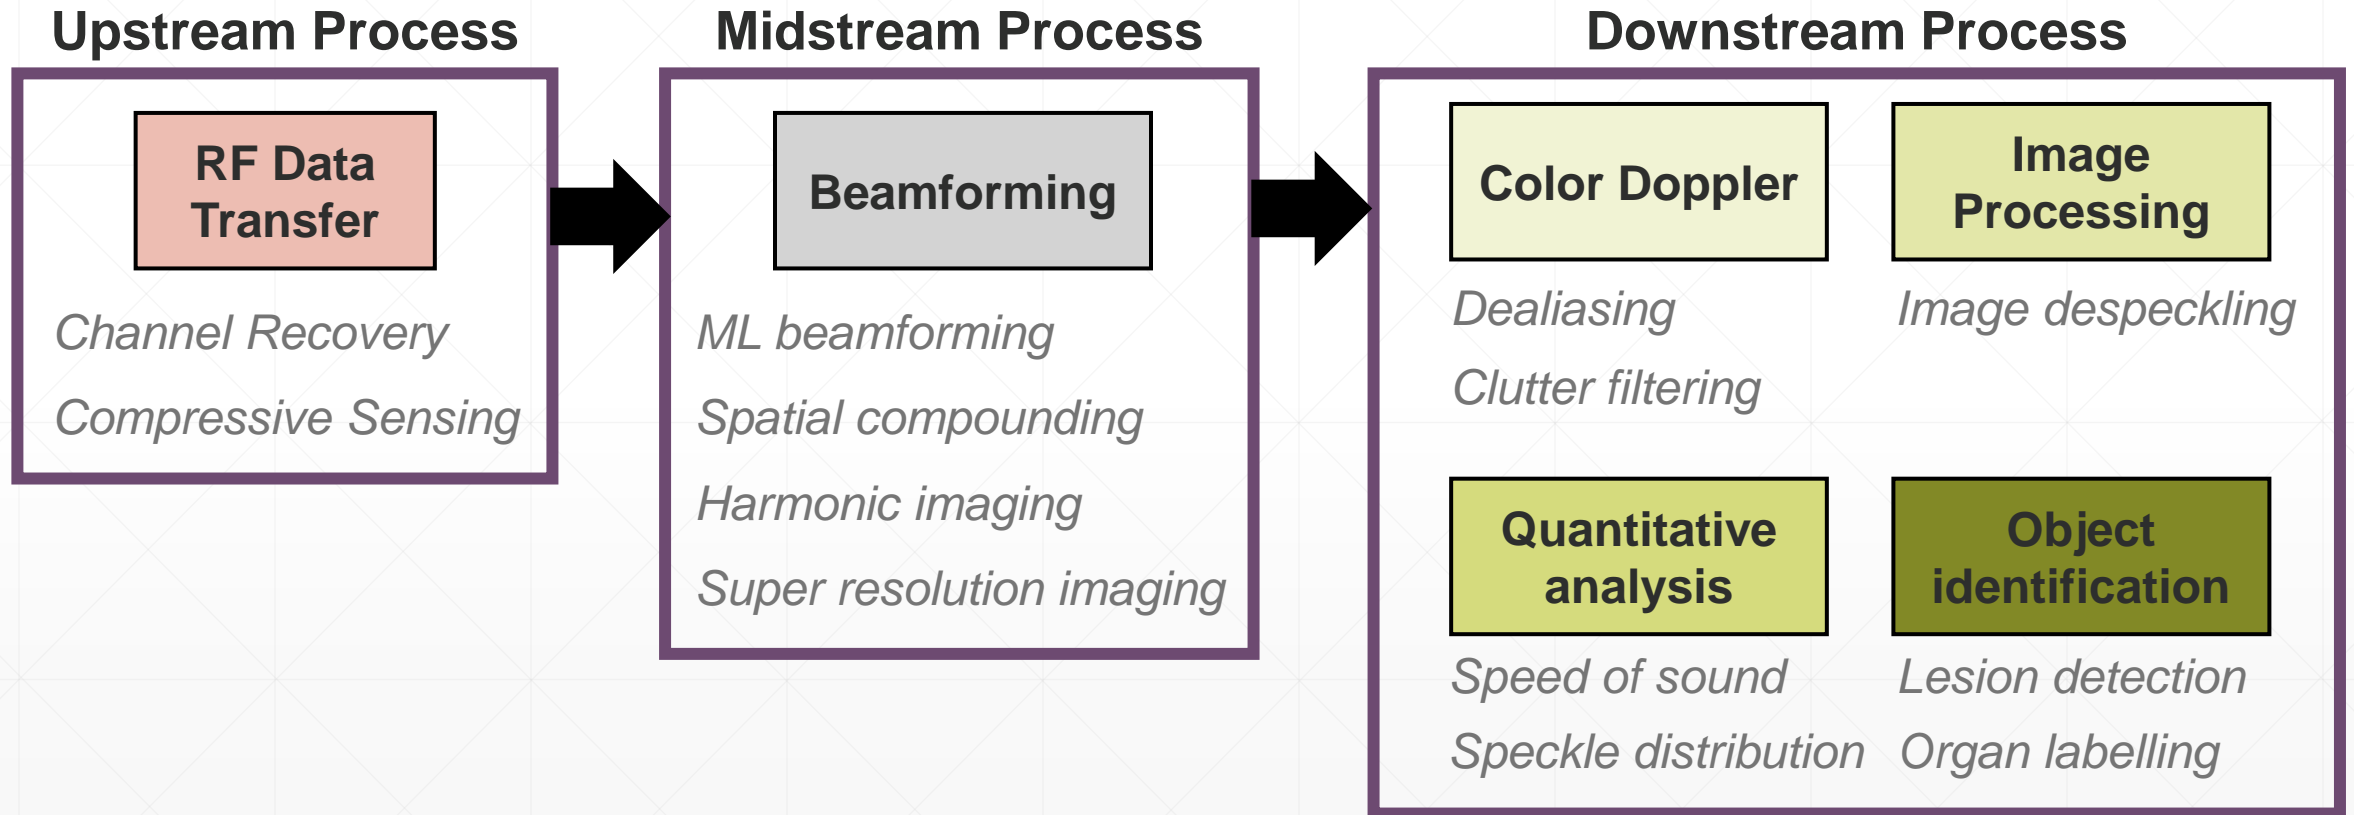
- Greater control on implementation
- Need to implement many things by yourselves
- Optimize the implementation specific to your algorithms for a real-time performance
- Libraries like cuBLAS, cuRAND, cuSOLVER, cuSPARSE, cuFFT, cuDNN are available
- Greater control on display rendering

C++ Deployment of Machine Learning Algorithms

Di Xiao, Hassan Nahas, Billy Yiu

Machine Learning in Ultrasound Imaging

Involved in every stage of ultrasound image processing



Deploying a Trained Neural Network in C++ Environment

Not a straightforward task

TensorFlow C++

TensorFlow API
directly in C++

Supports inference
on CPU and GPU

TensorRT

API for optimized inference
on NVIDIA devices.

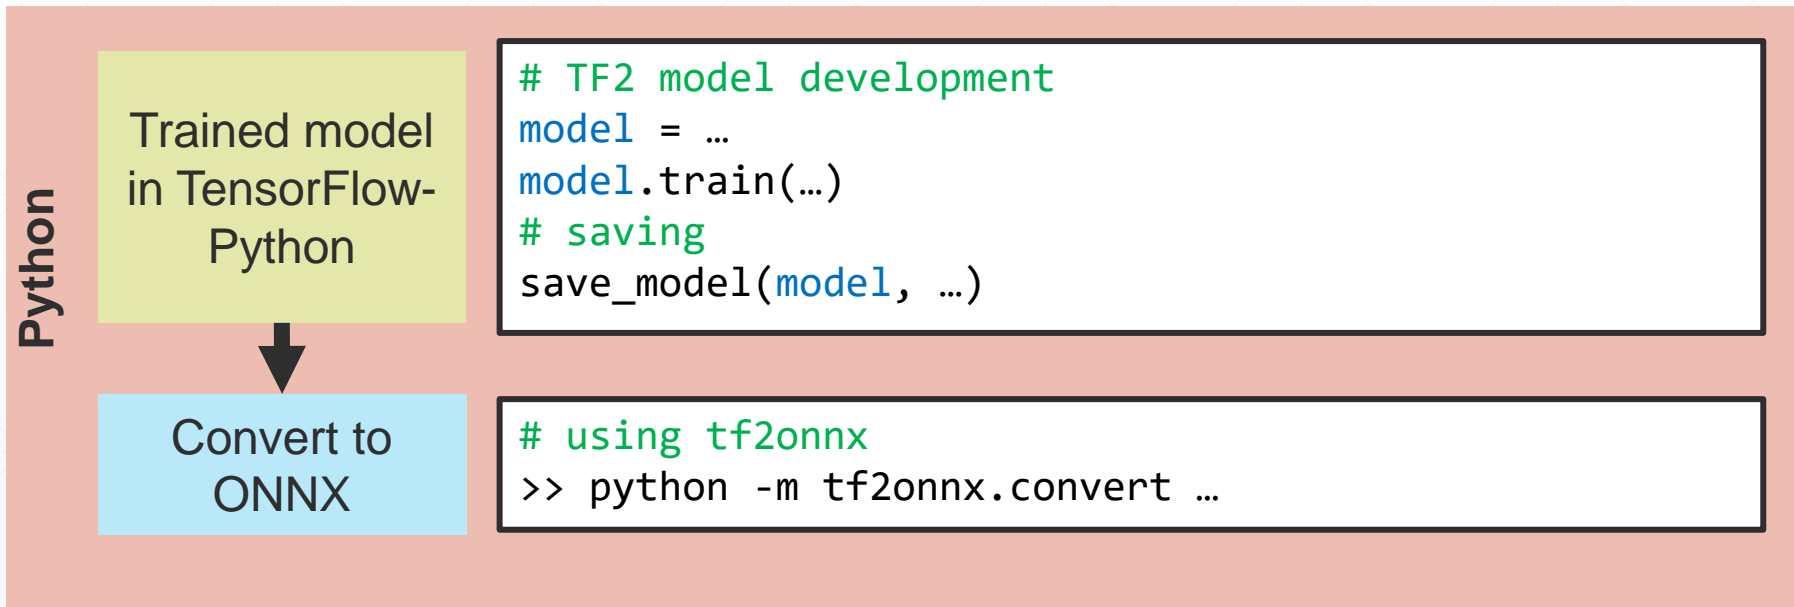

May be used with other
deep learning packages
(e.g.: Pytorch)

TensorFlow-TensorRT

Direct integration between TF and TRT

TensorRT for Model Deployment

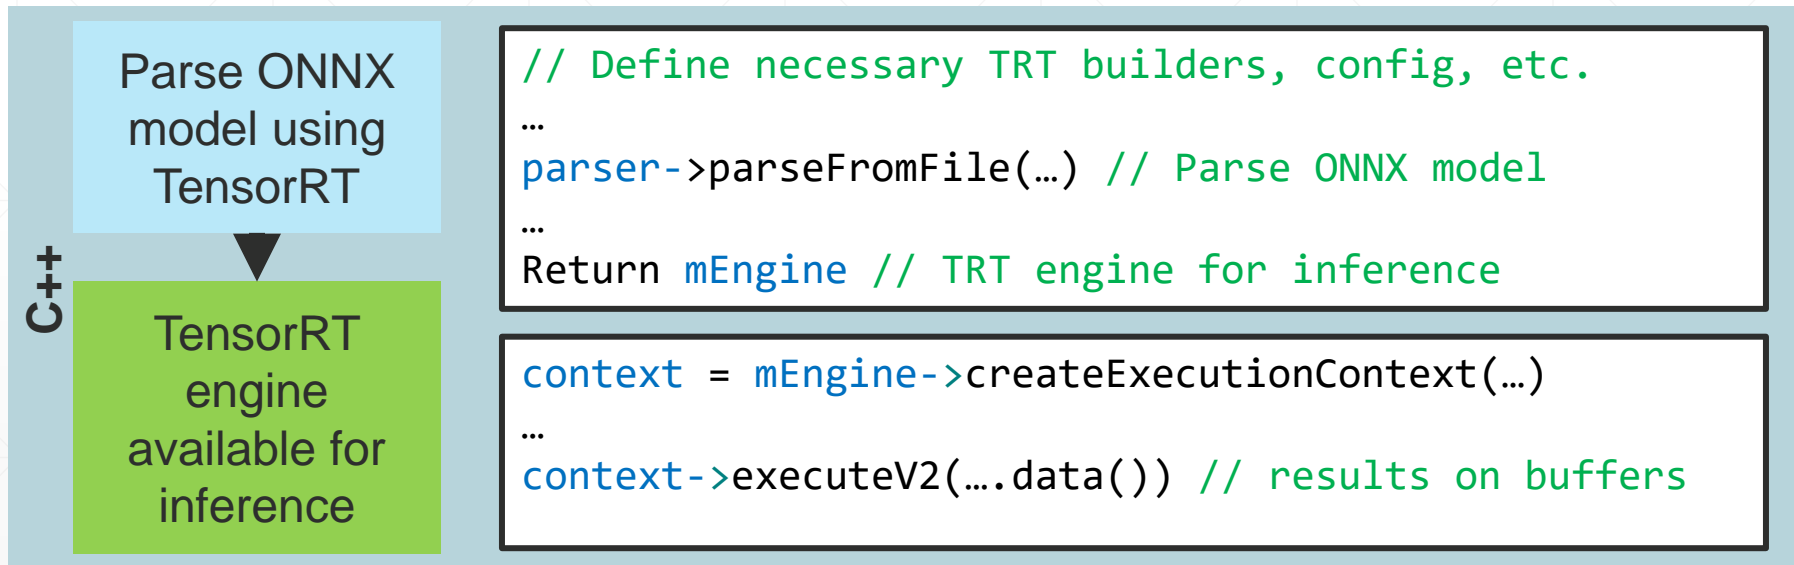

Save the model in TensorRT readable format: ONNX



TF2ONNX is open-source and available at <https://onnxruntime.ai/docs/tutorials/tf-get-started.html>

TensorRT for Model Deployment

Save the model in TensorRT readable format: ONNX



Deploying a CNN Model Trained with MNIST Dataset: A Brief Walkthrough

1. Define the TF-Keras model

```
#####  
# Creates the model.  
  
model = tf.keras.models.Sequential([  
    tf.keras.Input((28, 28, 1), batch_size=1),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(128, activation='relu'),  
    tf.keras.layers.Dropout(0.2),  
    tf.keras.layers.Dense(10),  
    tf.keras.layers.Softmax()  
])
```

2. Train the model with MNIST dataset

```
#####  
# Train the model on MNIST  
  
mnist = tf.keras.datasets.mnist  
  
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
x_train, x_test = x_train / 255.0, x_test / 255.0  
  
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False)  
  
model.compile(optimizer='adam',  
              loss=loss_fn,  
              metrics=['accuracy'])  
  
model.fit(x_train, y_train, epochs=5)  
  
model.evaluate(x_test, y_test, verbose=2)
```

*Based on sample codes by NVIDIA and ONNX project

Deploying a CNN Model Trained with MNIST Dataset: A Brief Walkthrough

3. Save the TensorFlow model

```
#####  
# Saves the model for ONNX conversion.  
  
if not os.path.exists("simple_nn"):  
    os.mkdir("simple_nn")  
tf.keras.models.save_model(model, "simple_nn")
```

4. Convert the saved TensorFlow model into an ONNX one

```
#####  
# Run the command line for ONNX conversion.  
  
proc = subprocess.run('python -m tf2onnx.convert --saved-model simple_nn '  
    '--output ../cpp/sample_mnist_data/simple_nn.onnx --opset 15'.split(),  
    capture_output=True)
```

*Based on sample codes by NVIDIA and ONNX project

Deploying a CNN Model Trained with MNIST Dataset: A Brief Walkthrough

5. Load the ONNX model into C++ environment

```
// create builder object
auto builder = SampleUniquePtr<nvinfer1::IBuilder>(nvinfer1::createInferBuilder(sample::gLogger.getTRTLogger()));

// create network object
const auto explicitBatch = 1U << static_cast<uint32_t>(NetworkDefinitionCreationFlag::kEXPLICIT_BATCH);
auto network = SampleUniquePtr<nvinfer1::INetworkDefinition>(builder->createNetworkV2(explicitBatch));

// create config object
auto config = SampleUniquePtr<nvinfer1::IBuilderConfig>(builder->createBuilderConfig());

// create parser object
auto parser = SampleUniquePtr<nvonnxparser::IParser>(nvonnxparser::createParser(*network, sample::gLogger.getTRTLogger()));

// parse ONNX model from filepath
auto parsed = parser->parseFromFile(onnxModelPath.c_str(), static_cast<int>(sample::gLogger.getReportableSeverity()));
```

*Based on sample codes by NVIDIA and ONNX project

Deploying a CNN Model Trained with MNIST Dataset: A Brief Walkthrough

6. Configure the CNN engine

```
// Configure DLA
samplesCommon::enableDLA(builder.get(), config.get(), dlaCore);

// Build serialized network and return it
SampleUniquePtr<IHostMemory> plan{ builder->buildSerializedNetwork(*network, *config) };

// Get inference runtime
mRuntime = std::shared_ptr<nvinfer1::IRuntime>(createInferRuntime(sample::gLogger.getTRTLogger()));

// Get inference engine for network
mEngine = std::shared_ptr<nvinfer1::ICudaEngine>(
    mRuntime->deserializeCudaEngine(plan->data(), plan->size()), samplesCommon::InferDeleter());
```

*Based on sample codes by NVIDIA and ONNX project

Deploying a CNN Model Trained with MNIST Dataset: A Brief Walkthrough

7. Execute the CNN for inferencing

```
// buffers for input/output of model
samplesCommon::BufferManager buffers(mEngine);

// Create execution context for inference from engine.
auto context = SampleUniquePtr<nvinfer1::IExecutionContext>(mEngine->createExecutionContext());

// Read the input data into the managed buffers
processInput(buffers);

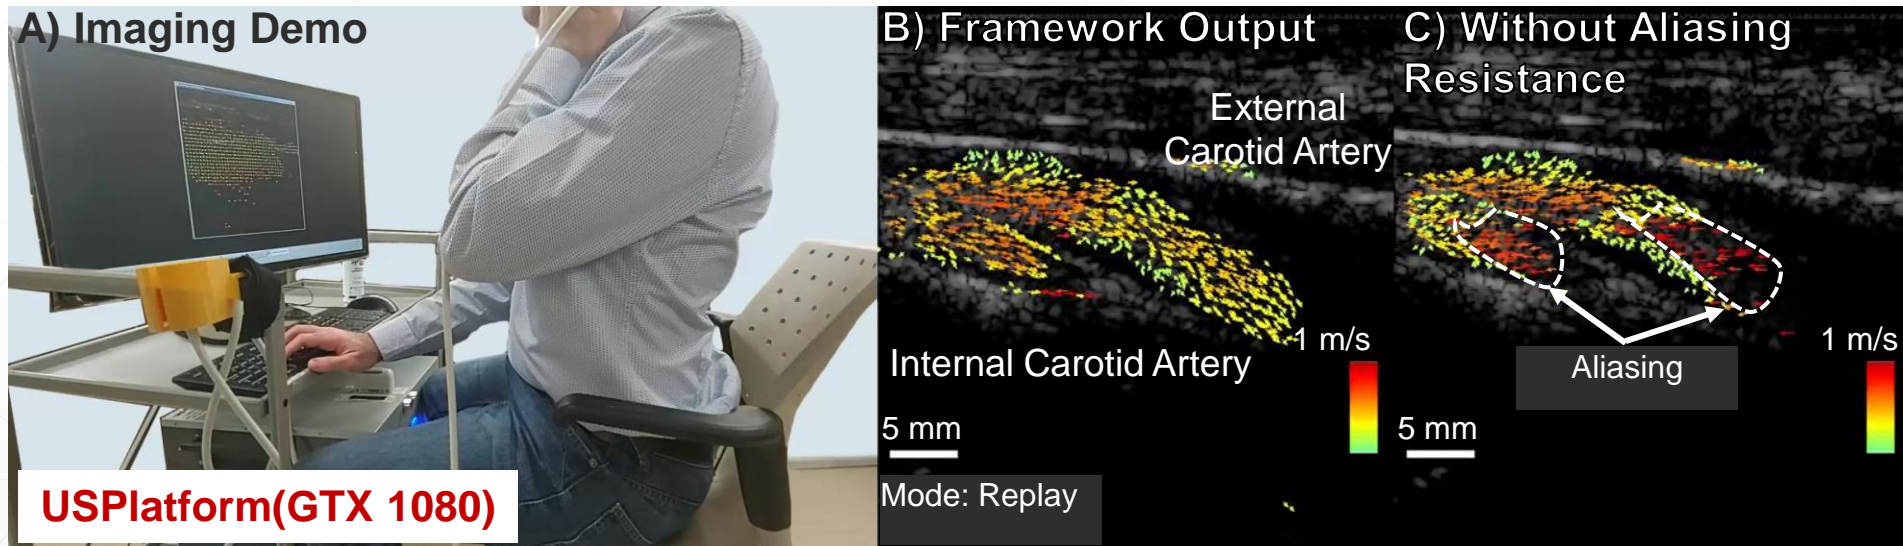
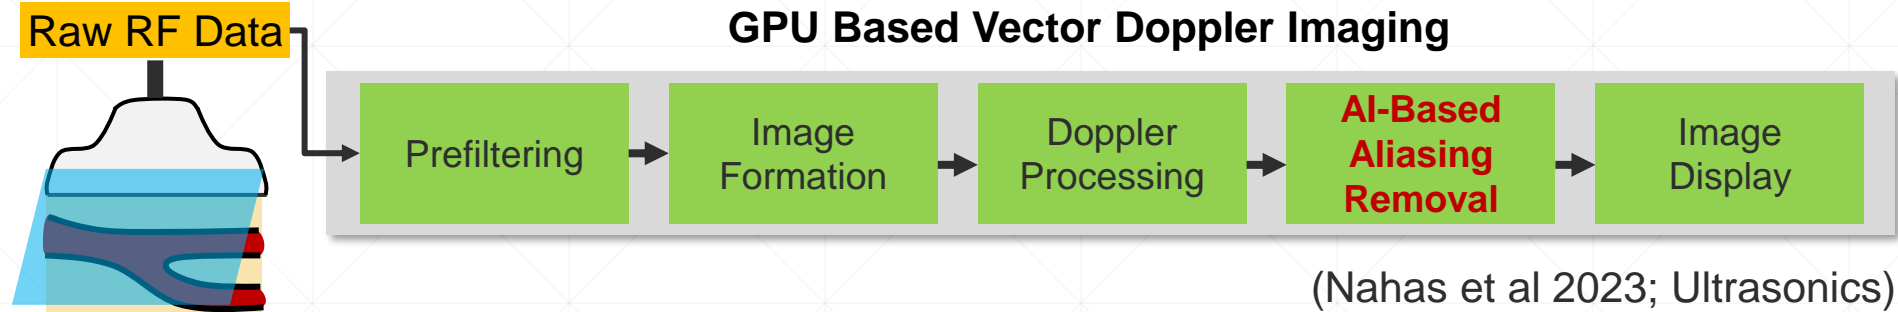
// Malloc from host input buffers to device input buffers
buffers.copyInputToDevice();

// Run inference once (for warming)
bool status = context->executeV2(buffers.getDeviceBindings().data()); // synronous inference on GPU

// Malloc from device output buffers to host output buffers
buffers.copyOutputToHost();
```

*Based on sample codes by NVIDIA and ONNX project

Demonstration: Real-time Aliasing Correction Using Deep Learning




Results:

- Live vector Doppler imaging at 16 fps used for guiding examination
- On-site aliasing-resistant playback at average 60 fps with a screen update rate of 30 fps

Short Summary on TensorRT Implementation

Use ONNX convertor to translate the network into a TensorRT readable format



ONNX

- Able to interface with many different packages
- Open source

TensorRT

- Able to leverage on the dedicated inference hardware on Nvidia GPU
- Make the NN readily available for existing processing pipeline